

hotel

March 26, 2021

0.1 Hotel Cancellations (Originally from Kaggle Courses)

We will build a model to predict hotel cancellations with a binary classifier.

```
[1]: import matplotlib.pyplot as plt
plt.style.use('seaborn-whitegrid')
plt.rc('figure', autolayout=True)
plt.rc('axes', labelweight='bold', labelsize='large',
       titleweight='bold', titlesize=18, titlepad=10)
plt.rc('animation', html='html5')
```

```
[2]: import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.compose import make_column_transformer

hotel = pd.read_csv('./hotel.csv')

X = hotel.copy() # Deep Copy
y = X.pop('is_canceled')
```

```
[3]: X.head()
```

```
[3]:      hotel  lead_time  arrival_date_year  arrival_date_month \
0  Resort Hotel      342          2015          July
1  Resort Hotel      737          2015          July
2  Resort Hotel       7          2015          July
3  Resort Hotel      13          2015          July
4  Resort Hotel      14          2015          July

      arrival_date_week_number  arrival_date_day_of_month \
0                27                1
1                27                1
2                27                1
3                27                1
```

```

4                                27                                1

    stays_in_weekend_nights  stays_in_week_nights  adults  children  ...  \
0                          0                      0       2        0.0  ...
1                          0                      0       2        0.0  ...
2                          0                      1       1        0.0  ...
3                          0                      1       1        0.0  ...
4                          0                      2       2        0.0  ...

    deposit_type  agent  company  days_in_waiting_list  customer_type  adr  \
0    No Deposit   NaN    NaN                      0    Transient    0.0
1    No Deposit   NaN    NaN                      0    Transient    0.0
2    No Deposit   NaN    NaN                      0    Transient   75.0
3    No Deposit  304.0    NaN                      0    Transient   75.0
4    No Deposit  240.0    NaN                      0    Transient   98.0

    required_car_parking_spaces  total_of_special_requests  reservation_status  \
0                              0                          0          Check-Out
1                              0                          0          Check-Out
2                              0                          0          Check-Out
3                              0                          0          Check-Out
4                              0                          1          Check-Out

    reservation_status_date
0          2015-07-01
1          2015-07-01
2          2015-07-02
3          2015-07-02
4          2015-07-03

```

[5 rows x 31 columns]

```
[4]: y.head()
```

```

[4]: 0    0
     1    0
     2    0
     3    0
     4    0
     Name: is_canceled, dtype: int64

```

```

[5]: # Mapping dates to numbers

X['arrival_date_month'] = X['arrival_date_month'].map({'January':1, 'February':1,
→2, 'March':3,
                                                    'April':4, 'May':5,
→'June':6, 'July':7,

```

```

        'August':8, 'September':
↪9, 'October':10,
        'November':11,
↪'December':12})

```

[6]: *# Defining features*

```

features_num = ["lead_time", "arrival_date_week_number",
                "arrival_date_day_of_month", "stays_in_weekend_nights",
                "stays_in_week_nights", "adults", "children", "babies",
                "is_repeated_guest", "previous_cancellations",
                "previous_bookings_not_canceled", "required_car_parking_spaces",
                "total_of_special_requests", "adr"]

```

[7]: *# Defining categorical features*

```

features_cat = ["hotel", "arrival_date_month", "meal",
                "market_segment", "distribution_channel",
                "reserved_room_type", "deposit_type", "customer_type"]

```

[8]: *# Transforming numerical and categorical features*

```

transformer_num = make_pipeline(SimpleImputer(strategy="constant"), # there are
↪a few missing values
                                StandardScaler())

transformer_cat = make_pipeline(SimpleImputer(strategy="constant",
↪fill_value="NA"),
                                OneHotEncoder(handle_unknown='ignore'))

preprocessor = make_column_transformer((transformer_num, features_num),
                                       (transformer_cat, features_cat))

```

[9]: *# stratify - make sure classes are evenly represented across splits*

```

X_train, X_valid, y_train, y_valid = train_test_split(X, y, stratify=y,
↪train_size=0.75)

X_train = preprocessor.fit_transform(X_train)
X_valid = preprocessor.transform(X_valid)

input_shape = [X_train.shape[1]]
print(input_shape)

```

[63]

0.2 Defining a Model

The model will have both batch normalization and dropout layers.

```
[10]: from tensorflow import keras
      from tensorflow.keras import layers

      model = keras.Sequential([layers.BatchNormalization(),
                                layers.Dense(256, activation='relu',
                                ↪input_shape=input_shape),
                                layers.BatchNormalization(),
                                layers.Dropout(rate=0.3),
                                layers.Dense(256, activation='relu'),
                                layers.BatchNormalization(),
                                layers.Dropout(rate=0.3),
                                layers.Dense(1, activation='sigmoid')])
```

0.3 Optimizer, Loss Function, and Accuracy Metric

We use Adam optimizer and binary versions of the cross-entropy loss and accuracy metric.

```
[11]: model.compile(optimizer='adam', loss='binary_crossentropy',
      ↪metrics=['binary_accuracy'])
```

0.4 Early stopping to prevent overfitting

```
[12]: early_stopping = keras.callbacks.EarlyStopping(patience=5,min_delta=0.
      ↪001,restore_best_weights=True)

      history = model.fit(X_train, y_train,
                          validation_data=(X_valid, y_valid),
                          batch_size=512,
                          epochs=200,
                          callbacks=[early_stopping])
```

Epoch 1/200

175/175 [=====] - 3s 10ms/step - loss: 0.5533 -
binary_accuracy: 0.7312 - val_loss: 0.4381 - val_binary_accuracy: 0.8044

Epoch 2/200

175/175 [=====] - 1s 8ms/step - loss: 0.4311 -
binary_accuracy: 0.7976 - val_loss: 0.4026 - val_binary_accuracy: 0.8138

Epoch 3/200

175/175 [=====] - 1s 8ms/step - loss: 0.4133 -
binary_accuracy: 0.8075 - val_loss: 0.3988 - val_binary_accuracy: 0.8138

Epoch 4/200

175/175 [=====] - 1s 8ms/step - loss: 0.4009 -
binary_accuracy: 0.8117 - val_loss: 0.3913 - val_binary_accuracy: 0.8175

Epoch 5/200


175/175 [=====] - 1s 8ms/step - loss: 0.3978 -
binary_accuracy: 0.8147 - val_loss: 0.3873 - val_binary_accuracy: 0.8209
Epoch 6/200
175/175 [=====] - 1s 8ms/step - loss: 0.3906 -
binary_accuracy: 0.8185 - val_loss: 0.3864 - val_binary_accuracy: 0.8213
Epoch 7/200
175/175 [=====] - 1s 8ms/step - loss: 0.3900 -
binary_accuracy: 0.8184 - val_loss: 0.3817 - val_binary_accuracy: 0.8247
Epoch 8/200
175/175 [=====] - 1s 8ms/step - loss: 0.3834 -
binary_accuracy: 0.8215 - val_loss: 0.3837 - val_binary_accuracy: 0.8239
Epoch 9/200
175/175 [=====] - 1s 8ms/step - loss: 0.3818 -
binary_accuracy: 0.8211 - val_loss: 0.3803 - val_binary_accuracy: 0.8265
Epoch 10/200
175/175 [=====] - 1s 8ms/step - loss: 0.3794 -
binary_accuracy: 0.8224 - val_loss: 0.3762 - val_binary_accuracy: 0.8300
Epoch 11/200
175/175 [=====] - 1s 8ms/step - loss: 0.3764 -
binary_accuracy: 0.8249 - val_loss: 0.3737 - val_binary_accuracy: 0.8328
Epoch 12/200
175/175 [=====] - 1s 8ms/step - loss: 0.3735 -
binary_accuracy: 0.8267 - val_loss: 0.3691 - val_binary_accuracy: 0.8318
Epoch 13/200
175/175 [=====] - 1s 8ms/step - loss: 0.3715 -
binary_accuracy: 0.8279 - val_loss: 0.3686 - val_binary_accuracy: 0.8322
Epoch 14/200
175/175 [=====] - 1s 8ms/step - loss: 0.3703 -
binary_accuracy: 0.8279 - val_loss: 0.3683 - val_binary_accuracy: 0.8324
Epoch 15/200
175/175 [=====] - 1s 8ms/step - loss: 0.3670 -
binary_accuracy: 0.8296 - val_loss: 0.3672 - val_binary_accuracy: 0.8328
Epoch 16/200
175/175 [=====] - 2s 9ms/step - loss: 0.3632 -
binary_accuracy: 0.8324 - val_loss: 0.3666 - val_binary_accuracy: 0.8347
Epoch 17/200
175/175 [=====] - 1s 8ms/step - loss: 0.3638 -
binary_accuracy: 0.8301 - val_loss: 0.3637 - val_binary_accuracy: 0.8346
Epoch 18/200
175/175 [=====] - 1s 8ms/step - loss: 0.3639 -
binary_accuracy: 0.8319 - val_loss: 0.3679 - val_binary_accuracy: 0.8341
Epoch 19/200
175/175 [=====] - 1s 8ms/step - loss: 0.3591 -
binary_accuracy: 0.8340 - val_loss: 0.3649 - val_binary_accuracy: 0.8345
Epoch 20/200
175/175 [=====] - 2s 9ms/step - loss: 0.3568 -
binary_accuracy: 0.8338 - val_loss: 0.3584 - val_binary_accuracy: 0.8375
Epoch 21/200


175/175 [=====] - 2s 9ms/step - loss: 0.3575 -
binary_accuracy: 0.8347 - val_loss: 0.3579 - val_binary_accuracy: 0.8373
Epoch 22/200
175/175 [=====] - 2s 10ms/step - loss: 0.3543 -
binary_accuracy: 0.8354 - val_loss: 0.3591 - val_binary_accuracy: 0.8382
Epoch 23/200
175/175 [=====] - 2s 10ms/step - loss: 0.3503 -
binary_accuracy: 0.8385 - val_loss: 0.3568 - val_binary_accuracy: 0.8390
Epoch 24/200
175/175 [=====] - 2s 9ms/step - loss: 0.3513 -
binary_accuracy: 0.8373 - val_loss: 0.3577 - val_binary_accuracy: 0.8381
Epoch 25/200
175/175 [=====] - 2s 10ms/step - loss: 0.3498 -
binary_accuracy: 0.8383 - val_loss: 0.3578 - val_binary_accuracy: 0.8390
Epoch 26/200
175/175 [=====] - 2s 10ms/step - loss: 0.3501 -
binary_accuracy: 0.8385 - val_loss: 0.3575 - val_binary_accuracy: 0.8390
Epoch 27/200
175/175 [=====] - 2s 10ms/step - loss: 0.3453 -
binary_accuracy: 0.8392 - val_loss: 0.3560 - val_binary_accuracy: 0.8403
Epoch 28/200
175/175 [=====] - 2s 10ms/step - loss: 0.3490 -
binary_accuracy: 0.8372 - val_loss: 0.3552 - val_binary_accuracy: 0.8410
Epoch 29/200
175/175 [=====] - 2s 10ms/step - loss: 0.3460 -
binary_accuracy: 0.8402 - val_loss: 0.3548 - val_binary_accuracy: 0.8388
Epoch 30/200
175/175 [=====] - 2s 10ms/step - loss: 0.3446 -
binary_accuracy: 0.8395 - val_loss: 0.3539 - val_binary_accuracy: 0.8387
Epoch 31/200
175/175 [=====] - 2s 10ms/step - loss: 0.3419 -
binary_accuracy: 0.8416 - val_loss: 0.3591 - val_binary_accuracy: 0.8392
Epoch 32/200
175/175 [=====] - 2s 10ms/step - loss: 0.3416 -
binary_accuracy: 0.8418 - val_loss: 0.3566 - val_binary_accuracy: 0.8397
Epoch 33/200
175/175 [=====] - 2s 10ms/step - loss: 0.3424 -
binary_accuracy: 0.8395 - val_loss: 0.3551 - val_binary_accuracy: 0.8388
Epoch 34/200
175/175 [=====] - 2s 10ms/step - loss: 0.3428 -
binary_accuracy: 0.8408 - val_loss: 0.3509 - val_binary_accuracy: 0.8424
Epoch 35/200
175/175 [=====] - 2s 11ms/step - loss: 0.3419 -
binary_accuracy: 0.8417 - val_loss: 0.3539 - val_binary_accuracy: 0.8408
Epoch 36/200
175/175 [=====] - 2s 11ms/step - loss: 0.3412 -
binary_accuracy: 0.8418 - val_loss: 0.3562 - val_binary_accuracy: 0.8399
Epoch 37/200

```
175/175 [=====] - 2s 10ms/step - loss: 0.3397 -  
binary_accuracy: 0.8425 - val_loss: 0.3523 - val_binary_accuracy: 0.8427  
Epoch 38/200  
175/175 [=====] - 2s 11ms/step - loss: 0.3399 -  
binary_accuracy: 0.8423 - val_loss: 0.3520 - val_binary_accuracy: 0.8412  
Epoch 39/200  
175/175 [=====] - 2s 11ms/step - loss: 0.3358 -  
binary_accuracy: 0.8450 - val_loss: 0.3502 - val_binary_accuracy: 0.8429
```

```
[13]: history_df = pd.DataFrame(history.history)  
history_df.loc[:, ['loss', 'val_loss']].plot(title="Cross-entropy")  
history_df.loc[:, ['binary_accuracy', 'val_binary_accuracy']].  
    ↪plot(title="Accuracy")
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f965e6e1f40>
```





[]: