

# 1.wine

January 17, 2021

```
[1]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn import svm
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
```

```
[2]: wine = pd.read_csv('./datasets/winequality-red.csv',sep=';')
```

```
[3]: wine.head()
```

```
[3]: fixed acidity,volatile acidity,citric acid,residual sugar,chlorides,free
sulfur dioxide,total sulfur dioxide,density,pH,sulphates,alcohol,quality
0  7.4,0.7,0.0,1.9,0.076,11.0,34.0,0.9978,3.51,0...
1  7.8,0.88,0.0,2.6,0.098,25.0,67.0,0.9968,3.2,0...
2  7.8,0.76,0.04,2.3,0.092,15.0,54.0,0.997,3.26,0...
3  11.2,0.28,0.56,1.9,0.075,17.0,60.0,0.998,3.16,...
4  7.4,0.7,0.0,1.9,0.076,11.0,34.0,0.9978,3.51,0...
```

```
[4]: wine = pd.read_csv('./datasets/winequality-red.csv',sep=',')
```

```
[5]: wine.head()
```

```
[5]: fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4           0.70           0.00           1.9           0.076
1           7.8           0.88           0.00           2.6           0.098
2           7.8           0.76           0.04           2.3           0.092
3          11.2           0.28           0.56           1.9           0.075
4           7.4           0.70           0.00           1.9           0.076

free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0              11.0              34.0  0.9978  3.51           0.56
1              25.0              67.0  0.9968  3.20           0.68
```

2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[6]: wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
[7]: wine.isnull().sum()
```

```
[7]: fixed acidity          0
     volatile acidity      0
     citric acid           0
     residual sugar        0
     chlorides             0
     free sulfur dioxide    0
     total sulfur dioxide   0
     density               0
     pH                   0
     sulphates             0
     alcohol               0
     quality               0
```

dtype: int64

```
[8]: #preprocessing Data
bins = (2,6.5,8)
group_names = ['bad', 'good']
wine['quality'] = pd.cut(wine['quality'], bins = bins, labels = group_names)
wine['quality'].unique()
```

```
[8]: [bad, good]
Categories (2, object): [bad < good]
```

```
[9]: label_quality = LabelEncoder()
```

```
[10]: wine['quality'] = label_quality.fit_transform(wine['quality'])
```

```
[11]: wine.head()
```

```
[11]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4             0.70         0.00           1.9       0.076
1           7.8             0.88         0.00           2.6       0.098
2           7.8             0.76         0.04           2.3       0.092
3          11.2             0.28         0.56           1.9       0.075
4           7.4             0.70         0.00           1.9       0.076
```

```
   free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0              11.0             34.0  0.9978  3.51       0.56
1              25.0             67.0  0.9968  3.20       0.68
2              15.0             54.0  0.9970  3.26       0.65
3              17.0             60.0  0.9980  3.16       0.58
4              11.0             34.0  0.9978  3.51       0.56
```

```
   alcohol  quality
0       9.4        0
1       9.8        0
2       9.8        0
3       9.8        0
4       9.4        0
```

```
[12]: wine.head(10)
```

```
[12]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4             0.70         0.00           1.9       0.076
1           7.8             0.88         0.00           2.6       0.098
2           7.8             0.76         0.04           2.3       0.092
3          11.2             0.28         0.56           1.9       0.075
4           7.4             0.70         0.00           1.9       0.076
5           7.4             0.66         0.00           1.8       0.075
```

6	7.9	0.60	0.06	1.6	0.069
7	7.3	0.65	0.00	1.2	0.065
8	7.8	0.58	0.02	2.0	0.073
9	7.5	0.50	0.36	6.1	0.071

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56
5	13.0	40.0	0.9978	3.51	0.56
6	15.0	59.0	0.9964	3.30	0.46
7	15.0	21.0	0.9946	3.39	0.47
8	9.0	18.0	0.9968	3.36	0.57
9	17.0	102.0	0.9978	3.35	0.80


	alcohol	quality
0	9.4	0
1	9.8	0
2	9.8	0
3	9.8	0
4	9.4	0
5	9.4	0
6	9.4	0
7	10.0	1
8	9.5	1
9	10.5	0

```
[13]: wine['quality'].value_counts()
```

```
[13]: 0    1382
      1     217
      Name: quality, dtype: int64
```

```
[14]: sns.countplot(wine['quality'])
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7fbeb436490>
```



```
[15]: #separate the dataset as response variable and feature variables
```

```
X = wine.drop('quality', axis=1)
y = wine['quality']
```

```
[16]: #Train and Test splitting of data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[17]: #Applying Standard scaling to get optimized result
```

```
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
[18]: X_train[:10]
```

```
[18]: array([[ 0.21833164,  0.88971201,  0.19209222,  0.30972563, -0.04964208,
  0.69100692,  1.04293362,  1.84669643,  1.09349989,  0.45822284,
  1.12317723],
 [-1.29016623, -1.78878251,  0.65275338, -0.80507963, -0.45521361,
  2.38847304,  3.59387025, -3.00449133, -0.40043872, -0.40119696,
  1.40827174],
 [ 1.49475291, -0.78434707,  1.01104539, -0.52637831,  0.59927236,
 -0.95796016, -0.99174203,  0.76865471, -0.07566946,  0.51551749,
 -0.58738978],
```

```
[ 0.27635078,  0.86181102, -0.06383064, -0.66572897, -0.00908493,
  0.01202048, -0.71842739,  0.08948842,  0.05423824, -1.08873281,
 -0.96751578],
[ 0.04427419,  2.81487994, -0.62686095,  2.39998549, -0.31326357,
 -0.47296984,  0.2229897 ,  1.1998714 ,  0.37900751, -0.9741435 ,
 -0.49235828],
[-0.07176411, -0.78434707,  1.11341454, -0.17800167,  0.21397941,
  3.01896045,  2.62208486,  0.60694845,  0.44396136,  1.89058918,
 -0.58738978],
[-1.17412793,  0.10848444, -0.62686095, -0.52637831, -0.23214927,
  0.98200112, -0.35400787, -1.95879086,  0.05423824,  0.91658007,
  1.12317723],
[-0.1878024 , -0.17052541,  0.60156881,  0.03102432, -0.13075639,
 -0.37597178, -0.01995665,  0.93036097,  0.76873063, -0.229313 ,
  0.26789373],
[-0.07176411,  0.61070216, -0.01264607, -0.38702766,  0.13286511,
 -1.05495822,  0.92146044,  0.37516948, -1.17988496, -0.229313 ,
 -1.25261029],
[ 1.8428678 , -1.95618842,  1.21578369,  1.00647892,  0.31537229,
 -1.15195628, -0.71842739,  1.52328391, -0.20557717,  1.77599987,
 -0.30229528]])
```

## 1 Random Forest Classifier

```
[21]: rfc = RandomForestClassifier(n_estimators=200)
      rfc.fit(X_train, y_train)
      pred_rfc = rfc.predict(X_test)
```

```
[22]: pred_rfc[:20]
```

```
[22]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0])
```

```
[24]: #how our model performed
      print(classification_report(y_test, pred_rfc))
      print(confusion_matrix(y_test, pred_rfc))
```

	precision	recall	f1-score	support
0	0.92	0.97	0.94	273
1	0.73	0.51	0.60	47
accuracy			0.90	320
macro avg	0.82	0.74	0.77	320
weighted avg	0.89	0.90	0.89	320

```
[[264  9]
```

```
[ 23  24]]
```

## 2 SVM Classifier

```
[25]: clf = svm.SVC()
      clf.fit(X_train, y_train)
      pred_clf = clf.predict(X_test)
```

```
[26]: #how our model performed
      print(classification_report(y_test, pred_clf))
      print(confusion_matrix(y_test, pred_clf))
```

	precision	recall	f1-score	support
0	0.88	0.98	0.93	273
1	0.71	0.26	0.37	47
accuracy			0.88	320
macro avg	0.80	0.62	0.65	320
weighted avg	0.86	0.88	0.85	320

```
[[268  5]
 [ 35 12]]
```

## 3 Neural Network

```
[39]: mlpc = MLPClassifier(hidden_layer_sizes=(11,11,11), max_iter=1000)
      mlpc.fit(X_train, y_train)
      pred_mlpc = mlpc.predict(X_test)
```

```
[40]: #how our model performed
      print(classification_report(y_test, pred_mlpc))
      print(confusion_matrix(y_test, pred_mlpc))
```

	precision	recall	f1-score	support
0	0.92	0.92	0.92	273
1	0.55	0.55	0.55	47
accuracy			0.87	320
macro avg	0.74	0.74	0.74	320
weighted avg	0.87	0.87	0.87	320

```
[[252  21]
 [ 21  26]]
```

```
[41]: from sklearn.metrics import accuracy_score
cm = accuracy_score(y_test, pred_rfc)
cm
```

[41]: 0.9

```
[42]: wine.head(10)
```

```
[42]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
5	7.4	0.66	0.00	1.8	0.075	
6	7.9	0.60	0.06	1.6	0.069	
7	7.3	0.65	0.00	1.2	0.065	
8	7.8	0.58	0.02	2.0	0.073	
9	7.5	0.50	0.36	6.1	0.071	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	
5	13.0	40.0	0.9978	3.51	0.56	
6	15.0	59.0	0.9964	3.30	0.46	
7	15.0	21.0	0.9946	3.39	0.47	
8	9.0	18.0	0.9968	3.36	0.57	
9	17.0	102.0	0.9978	3.35	0.80	

	alcohol	quality
0	9.4	0
1	9.8	0
2	9.8	0
3	9.8	0
4	9.4	0
5	9.4	0
6	9.4	0
7	10.0	1
8	9.5	1
9	10.5	0

```
[45]: #check a new wine with new features to be a good or bad wine
Xnew = [[7.3,0.58,0.00,2.0,0.065,15.0,21.0,0.9946,3.36,0.47,10.0]]
Xnew = sc.transform(Xnew)
```



```
ynew = rfc.predict(Xnew)  
ynew
```

[45]: array([0])

[ ]: